

# Filetypes for Annotation & Alignment

J Fass | 20 June 2017

# Filetypes

- Fasta
- Fastq
- GTF / GFF
- SAM / BAM / CRAM
- BED

# Fasta

```
>sequenceName | plus other junk | few maintain a standard here
```

AGTGGAGCAAGCACAGAGAAGAACTGCAGTCAGGACATAAAGTAAAGTA

ATTAATCTAAAAAATAGTCTGAGCAGTCTTCTCTGCTGAANNNNNNNNNN Assembly gap?

NNNNNNNNNNNNNNNNNNNNNAAATTTCTTACTAGGAGGTCTTTAGTACAGA

TTCCTGATATGTAATTAATCACTAAATGTCTTTAATGGGATCTCTTTCTA

TTGAGATATTTGTAACTTTCTTCATGTGATTGGTTTACAGATATTCAGG

TTTCTGCAAATGGGTGCTGTCTATATTATAGAATTTTTAGTTGAAATTTT

CAAAATACTCTTTGagtattctcttgtaattataattactttacaagggttt

gtggggcatctctttcatttgatgattacatggttgcagtattctttttgt

tcttagtcagactgtataattgtctgtgaagtccagtaaacttttgaaag

Soft-masked repetitive  
sequence? Low  
confidence assembly?

# Fastq

[illegible]

```
@header
<sequence>
+(sometimes header?)
<base qualities>
```

# Fastq

[illegible]

Oct	Dec	Hex	Char	Oct	Dec	Hex	Char
000	0	00	NUL '\0' (null character)	100	64	40	@
001	1	01	SOH (start of heading)	101	65	41	A
002	2	02	STX (start of text)	102	66	42	B
003	3	03	ETX (end of text)	103	67	43	C
004	4	04	EOT (end of transmission)	104	68	44	D
005	5	05	ENQ (enquiry)	105	69	45	E
006	6	06	ACK (acknowledge)	106	70	46	F
007	7	07	BEL '\a' (bell)	107	71	47	G
010	8	08	BS '\b' (backspace)	110	72	48	H
011	9	09	HT '\t' (horizontal tab)	111	73	49	I
012	10	0A	LF '\n' (new line)	112	74	4A	J
013	11	0B	VT '\v' (vertical tab)	113	75	4B	K

# Fastq

Oct	Dec	Hex	Char
112	74	4A	J

$$74 - 33 = 41$$

Probability of error =  $10^{(-41 / 10)} \sim 0.0001$

41 is the “phred-scaled Q-value”

Standard FASTQ encodes qualities using “phred + 33” quality characters. See [https://en.wikipedia.org/wiki/FASTQ\\_format](https://en.wikipedia.org/wiki/FASTQ_format) for a good graphic about current and older encodings.

Common QC question: “how many reads have average of at least Q30?”

# SAM / BAM / CRAM!

<http://www.htslib.org/>

See also samtools man page: <http://samtools.sourceforge.net/>

SAM spec grew out of 1000 Genomes Project (see Li et al. 2009 *Bioinformatics* 25:2078)

SAM is plain text; BAM is binary, compressed version of SAM; CRAM is further compressed but not widely used / recognizable by many tools.

# SAM

$$[\dots]$$

@SQ SN:ctg103993 LN:217

@SQ SN:ctg103994 LN:222

@SQ SN:ctg103995 LN:205

@SQ SN:ctg103996 LN:210

```
@PG ID:bwa PN:bwa VN:0.7.13-r1126 CL:bwa mem -t 4 -M ../01_Reference/Transcriptome-Contigs-Build2.fna
```

```
../..02-Cleaned/3E/3E SE.fastq
```

@PG ID:bwa-7BC92A6F PN:bwa VN:0.7.13-r1126 CL:bwa mem -t 4 -M .././01 Reference:Transcriptome Coding Build2.1

```
../02-Cleaned/3E/3E R1.fastq ../02-Cleaned/3E/3E R2.fastq
```

K00188:264:HG3WJBBXX:1:1116:14692:35180#0	121	ctg2	128	58	101M =	128	0
---	-----	------	-----	----	--------	-----	---

AAGTCTCGACCAAGTGGTTCAGATGGTGACACAGATGTTAGCCCCATCCACCATTGAGTTGCCGTTTTGATAGCTGGAAATCCTGTAAACACAA

[illegible]

```
K00188:264:HG3WJBBXX:1:1116:14692:35180#0    181    ctg2    128    0    *    =    128    0
```

TTAGTTTTAATTTTGACTTTGAATAGCGGGAGTCCAGATCGTGTGAACACAGCAGACTGAGCACTCCATTGACAGCCTTCTTCTGTACTTTAGC

TATCC FJFJJFAAJF7F7JJJJAFFFAF&lt;7&lt;AFFJJJFJJJJJJJJJJJJJJJJJJJJFJJJJJJFAJJJJJJJJFFFJJJJJJJJJJFFJJJJJJJJFFAA AS:i:0 XS:i:0

K00188:264:HG3WJBBXX:1:1202:11028:9596#0	121	ctg5	45	60	101M	=	45	0
--	-----	------	----	----	------	---	----	---

TTCTTTTTCTACAGTTCATTGTCTGTATAAAGTATGCATCAGGAACAATCTGACTAGGAAGGTAAATAATGTAAACAGATGATTATTGTATGAAA

[illegible]

K00188:264:HG3WJBBXX:1:1202:11028:9596#0	181	ctg5	45	0	*	=	45	0
--	-----	------	----	---	---	---	----	---

TCAGCTGTATTAGTAATTTAGTAGAAAAGGTCTTGAGAGAATTATGTTTTTTAAAAATCCACATCACTTCAAACAAAAAGCCCCATTAGAATGGAGG

[illegible]

[...]

Header lines (start with “@”)

## Alignment line (one line per alignment)



# SAM

$$\begin{bmatrix} \vdots \\ \vdots \\ \vdots \end{bmatrix}$$

@SQ SN:ctg103993 LN:217

@SQ SN:ctg103994 LN:222

@SQ SN:ctg103995 LN:205

@SQ SN:ctg103996 LN:210

```
@PG ID:bwa PN:bwa VN:0.7.13-r1126 CL:bwa mem -t 4 -M ../01_Reference/Transcriptome-Contigs-Build2.fna
```

../..02-Cleaned/3E/3E SE.fastq

@PG ID:bwa-7BC92A6F PN:bwa VN:0.7.13-r1126 CL:bwa mem -t 4 -M .././01\_Reference/Transcriptome/ColinusBurditt/Read Pairs/Read Pair (identical headers)

```
../02-Cleaned/3E/3E R1.fastq ../02-Cleaned/3E/3E R2.fastq
```

K00188:264:HG3WJBBXX:1:1116:14692:35180#0	121	ctg2	128	58	101M =	128	0
---	-----	------	-----	----	--------	-----	---

AAGTCTCGACCAAGTGGTTCAGATGGTGACACAGATGTTAGCCCCATCCACCATTGAGTTGCCGTTTTGATAGCTGGAAATCCTGTAAACACAA

K00188:264:HG3WJBBXX:1:1116:14692:35180#0	181	ctg2	128	0	*	=	128	0
---	-----	------	-----	---	---	---	-----	---

TTAGTTTTTAATTTTGACTTTGAATAGCGGGAGTCCAGATCGTGTGAACACAGCAGACTGAGCACTCCATTGACAGCCTTCTTCTGTACTTTAGC

[illegible]

K00188:264:HG3WJBBXX:1:1202:11028:9596#0	121	ctg5	45	60	101M	=	45	0
--	-----	------	----	----	------	---	----	---

TTCTTTTTCTACAGTTCATTGTCTGTATAAAGTATGCATCAGGAACAATCTGACTAGGAAGGTAAATAATGTAAAACAGATGATTATTGTATGAAA

[illegible]

K00188:264:HG3WJBBXX:1:1202:11028:9596#0 181 ctg5 45 0 \* = 45 0

TCAGCTGTATTAGTAATTTAGTAGAAAAAGGTCTTGAGAGAAATTATGTTTTTTAAAAATCCACATCACTTCAAACAAAAAGCCCCATTAGAATGGAGG

[illegible]

[ 1 ]

[illegible]

Davis Genome Center | Bioinformatics Core | J Fass Filetypes 2017-06-20

# SAM

## QNAME: Query name

Read IDs are truncated at first whitespace (spaces / tabs), which can make them *non-unique*. Illumina reads with older IDs have trailing “/1” and “/2” stripped (this information is recorded in the next field). Illumina reads with newer IDs have second block stripped (read number is recorded in the next field).

@FCC6889ACXX:5:1101:8446:45501#CGATGTATC/1 ⇒ @FCC6889ACXX:5:1101:8446:45501

@HISEQ:153:H8ED7ADXX:1:1101:1368:2069 1:N:0:ATCACG ⇒ @HISEQ:153:H8ED7ADXX:1:1101:1368:2069

**HISEQ:153:H8ED7ADXX:1:1104:8193:69947**

99

chr1

4773690

50

101M

=

4773721

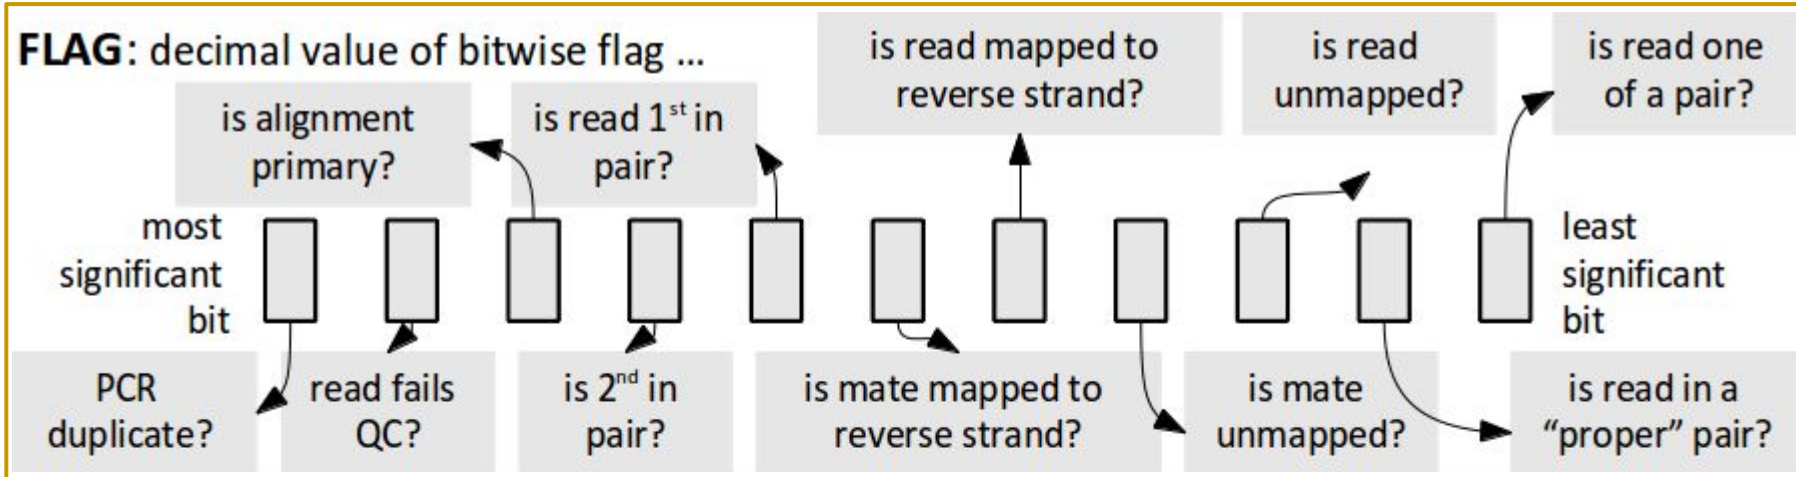
132

GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

**FLAG:** decimal value of bitwise flag ...



HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1  
4773690

50  
101M  
=  
4773721

132

GTGCCATCTGTGGGCTGGTGATC[... ]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[... ]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

99 (decimal) = 00001100011 (binary) ( 0 / NO .. 1 / YES )

... so, (from right to left): read is in a pair; the pair is proper; read is mapped (double neg); mate is mapped (double neg); read is mapped to forward strand (double neg); mate is mapped to reverse strand; read is 1st in pair ... *remaining bits not used*

# SAM

**FLAG:** still confused?

<https://broadinstitute.github.io/picard/explain-flags.html>

Common flags for SR (single reads): 0, 4, 16, sometimes 20 (hmm..)

Common flags for PE (paired ends): 99/147, 83/163, 77/141, 65/129, 81/161 ...

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**RNAME:** reference sequence name

Reference sequence ID (from fasta header), *possibly truncated at first whitespace (still unique??)*

>chromosome 1  
... *becomes* ...  
chromosome  
... (!)

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[... ]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[... ]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**POS:** 1-based *leftmost* position of (post-clipping) aligned read

... 4,773,680 4,773,690 4,773,700 4,773,710 ...

| | | |

REF: ... TACCCAATGGGGATGACATAAGGTGCCATCTGTGGGCTGGTGATTCCATAGTAGAC ...

READ: GGTGCCATCTGTGGGCTGGTGATCCCATAGTAGAC ...

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**POS:** 1-based *leftmost* position of (post-clipping) aligned read

... 4,773,680 4,773,690 4,773,700 4,773,710 ...

| | | |

REF: ... TACCCAATGGGGATGACATAAGGTGCCATCTGTGGGCTGGTGATTCCATAGTAGAC ...

READ: GGTGCCATCTGTGGGCTGGTGATCCCATAGTAGAC ...

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

49H101M

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1



# SAM

**POS:** 1-based *leftmost* position of (post-clipping) aligned read

... 4,773,680 4,773,690 4,773,700 4,773,710 ...

| | | |

REF: ... TACCCAATGGGGATGACATAAGGTGCCATCTGTGGGCTGGTGATTCCATAGTAGAC ...

READ: GCCGGTGCCATCTGTGGGCTGGTGATCCCATAGTAGAC ...

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

3S101M

=

4773721

132

GCCGTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBBBBBFFFFFBFFFFBFFFFBFFB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:GCC101(?) YT:Z:UU XS:A:- NH:i:1

# SAM

## MAPQ: mapping quality (phred scaled)

Mapping quality is used by some aligners, in different ways. It's generally a function of the edit distance (mismatches, indels), and the uniqueness of the alignment. Multiple equivalent best alignments yield a mapping quality of zero; alignments with an edit distance close to the best alignment lower the mapping quality.

```
HISEQ:153:H8ED7ADXX:1:1104:8193:69947
```

```
99
```

```
chr1
```

```
4773690
```

```
50
```

```
101M
```

```
=
```

```
4773721
```

```
132
```

```
GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG
```

```
BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBB<
```

```
AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1
```

# SAM

**CIGAR:** extended CIGAR string (Compact Idiosyncratic Gapped Alignment Report)

Format: [0-9][MIDNSHP][0-9][MIDNSHP]...

M = match / mismatch (!), I/D = insertion / deletion, N = skipped bases on reference, S/H = soft / hard clip (hard clipped bases no longer appear in the sequence field), P = padding

... e.g. "101M" means that all bases in the read align to bases in the reference, starting with position (4,773,690), always in the order of the reference.

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

**101M**

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[... ]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[... ]FFBFFBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**MRNM:** reference sequence to which the *mate* of this read is aligned

“=” ... mate is aligned to the same reference sequence as this read

“\*” ... this is a single read; no mate exists

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132

GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**MPOS:** 1-based, left-most position of 1st (post-clipping) nucleotide of mate read

“0” ... no mate exists

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132


GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1

# SAM

**TLEN:** inferred insert size / template length ... “0” if no mate ... “-#” if second read(?)



```
HISEQ:153:H8ED7ADXX:1:1104:8193:69947
```

```
99
```

```
chr1
```

```
4773690
```

```
50
```

```
101M
```

```
=
```

```
4773721
```

```
132
```

```
GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG
```

```
BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<
```

```
AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1
```

# SAM

**SEQ** and **QUAL**: read's nucleotides and base qualities, *always in the order of the reference (forward, top) strand!* ... includes any insertions, deletions, etc. present in the read.

Reads aligned to reverse strand appear in reverse, with reversed base qualities.

```
HISEQ:153:H8ED7ADXX:1:1104:8193:69947
```

```
99
```

```
chr1
```

```
4773690
```

```
50
```

```
101M
```

```
=
```

```
4773721
```

```
132
```

```
GTGCCATCTGTGGGCTGGTGATC[...]AGCAGCATGCTCCATGGTCTCTACATG
```

```
BBBFFFFFFBFFFFFFBFFBFFBB[...]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<
```

```
AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1
```

# SAM

**OPT:** various pre-defined and user-defined tags in the format TAG:VTYPE:VALUE ...  
VTYPE is one of [A (printable character); i (signed integer); f (floating point); z (printable string); H (hex string)].

e.g.: NM:i:0 means zero mismatches in this alignment

e.g.: XS:A:- was set by TopHat, RNA that was read was coded by the reverse strand

e.g.: NH:i:1 means that the number of hits for this read was 1 (would be more for repeat)

HISEQ:153:H8ED7ADXX:1:1104:8193:69947

99

chr1

4773690

50

101M

=

4773721

132


GTGCCATCTGTGGGCTGGTGATC[... ]AGCAGCATGCTCCATGGTCTCTACATG

BBBFFFFFFFFBFFFFFFFFBFFBFFBB[... ]FFBFFBBBBBBBBBBBBBBBBBBBBBBB<

AS:i:0 XN:i:0 XM:i:0 XO:i:0 XG:i:0 NM:i:0 MD:Z:101 YT:Z:UU XS:A:- NH:i:1



# SAM - quick summary



Ins & padding

Soft clipping

Splicing

Hard clipping

@SQ SN:ref LN:45

```
r001 163 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTA *
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003 0 ref 9 30 5H6M * 0 0 AGCTAA * NM:i:1
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC *
r003 16 ref 29 30 6H5M * 0 0 TAGGC * NM:i:0
r001 83 ref 37 30 9M = 7 -39 CAGCGCCAT *
```

ref 7 T 1 .	ref 12 T 3 ...	ref 17 T 3 ...
ref 8 T 1 .	ref 13 A 3 ...	ref 18 A 3 .-1G..
ref 9 A 3 ...	ref 14 A 2 .+2AG.+1G.	ref 19 G 2 *.
ref 10 G 3 ...	ref 15 G 2 ..	ref 20 C 2 ..
ref 11 A 3 ..C	ref 16 A 3 ...	...

google "Heng Li slides" - Challenges and Solutions in the Analysis of Next Generation Sequencing Data (2010)

# BAM

BAMs are compressed SAMs (so, binary, not human-readable text ... don't look directly at them!). They can be indexed to allow rapid extraction of information, so alignment viewers do not need to uncompress the whole BAM file in order to look at information for a particular read or coordinate range, somewhere in the file.

Indexing your BAM file, myCoolBamFile.bam, will create an index file, myCoolBamFile.bam.bai, which is needed (in addition to the BAM file) by viewers and other downstream tools. An occasional downstream tool will require an index called myCoolBamFile.bai (notice that the “.bai” replaces the “.bam”, instead of being appended after it).

# CRAM

Available as of SAMtools 1.0, and is a binary format like BAM. Uses data-specific compression tools (i.e. compressing letters is different than compressing numbers), *specifically* reference-based compression (e.g. for aligned reads, only *mis-matching* bases need to be stored). Also can employ *lossy* compression of base qualities, which appears to have a negligible effect on, say, variant calling (see Illumina [\*white paper\*](#)).

Indexing your CRAM file, myCoolBamFile.cram, will create an index file, myCoolBamFile.cram.crai, which is needed (in addition to the CRAM file) by viewers and other downstream tools.

This is still a ***relatively recent development***, so it may be a while before many tools are CRAM-capable.